

Sparse Voxel Octree

Tristano Munini

June 23, 2021

Topics and Objectives

The topics we will talk about are:

- ▶ Ray Tracing
- ▶ Signed Distance Functions
- ▶ Constructive Solid Geometry
- ▶ Sparse Voxel Octree (SVO)
- ▶ SVO Traversal Algorithm
- ▶ Implementation and Results

Objectives:


- ▶ use an SVO data structure to speed up the rendering process of my C++ implicit surface renderer
- ▶ get a better understanding of 3D spatial data structures

Ray Tracing: as simple as possible

- ▶ Can be seen as a root finding algorithm between rays and objects' descriptions
- ▶ Comprehends always a camera and a scene
- ▶ The rendering process generates an image out of a (3D) scene from the camera's perspective
- ▶ At each pixel corresponds at least one ray (usually several)
- ▶ The rays can bounce on surfaces, split in different rays and collect various information on their path
- ▶ Properly done Ray Tracing is extremely complex and computationally heavy

Ray Tracing: as simple as possible

© www.scratchapixel.com




In formulas:

$$r(t) = o + t * d$$

where t is a timestep, o is the ray's origin, d is the ray's direction, and $r(t)$ is a point along the ray.

Ray Tracing: as simple as possible



© www.scratchapixel.com

Signed Distance Functions

A distance surface is implicitly defined by a function $f: \mathbb{R}^3 \rightarrow \mathbb{R}$ that characterizes $A \subset \mathbb{R}^3$, set of points that are on or inside the implicit surface:

$$A = \{x : f(x) \leq 0\}$$

The surface can be also defined with $f^{-1}(0)$, which gives exactly the points on the surface.

Signed Distance Functions

The surface can be defined from the outside using a *point-to-set* distance:

$$d(x, A) = \min_{y \in A} \|x - y\|$$

Thus $d(x, A)$, given a point $x \in \mathbb{R}^3$, returns the shortest distance to the surface defined by A .

Definition (Signed Distance Function)

We say that f is a signed distance function (SDF) when holds

$$|f(x)| = d(x, f^{-1}(0)) \tag{1}$$

SDF Examples

Given a point $P = (x, y, z)$

- ▶ Sphere

$$\sqrt{x^2 + y^2 + z^2} - r$$

- ▶ Torus

$$\sqrt{\left(\sqrt{x^2 + z^2} - r_0\right)^2 + y^2} - r_1$$

- ▶ Cube

$$\sqrt{\max(|x| - l, 0)^2 + \max(|y| - l, 0)^2 + \max(|z| - l, 0)^2}$$


- ▶ Look the bibliography for more

Constructive Solid Geometry

SDFs make easy to create complex shapes from few simple primitives. This technique it known as Constructive Solid Geometry (CSG).

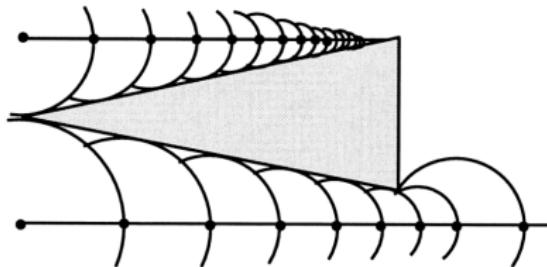
- ▶ union $\min(f_1, f_2)$
- ▶ intersection $\max(f_1, f_2)$
- ▶ subtraction $\max(f_1, -f_2)$
- ▶ mixing $k * f_1 + (1 - k) * f_2$ with $k \in [0, 1]$

CSG Example



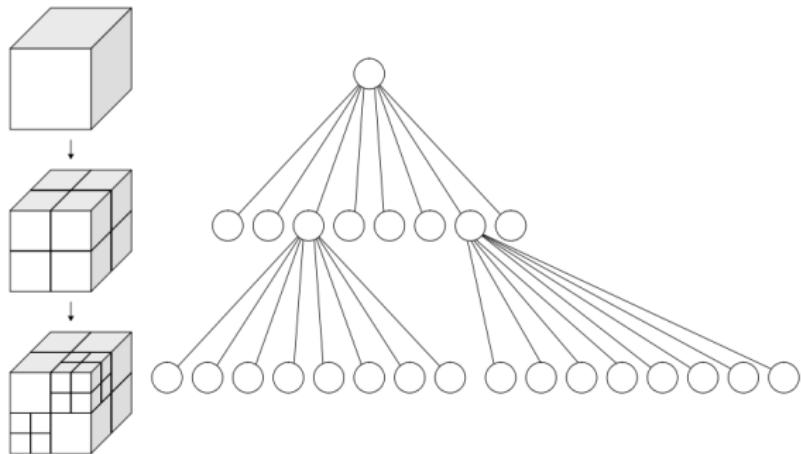
Sphere Tracing

- ▶ Studying ray's path proceeding with a fixed step may lead to errors, with Sphere Tracing an adaptive step is used
- ▶ The step size is given by an SDF
- ▶ Given $x \in \mathbb{R}^3$ and a surface S , $d(x, S)$ means we can move x in every direction by $d(x, S)$ being sure at worst to just hit the surface
- ▶ In other words we are drawing a safe sphere around the point, in which there is no surface




Octree

An Octree is a space partitioning data structure in which each node has exactly 8 children




Child Encoding

In a right-handed coordinate system the child numbering respects the quadrants' sing as if the cube was centered in the origin




Sparse Voxel Octree

- ▶ Complete information only in the leaves
 - ▶ Don't waste space encoding empty nodes



SVO Traversal

- ▶ Hierarchy traversal in depth-first order
- ▶ PUSH, ADVANCE, POP



SVO Traversal

- ▶ Each node is defined using its *parent* and an *idx* from 0 to 7
- ▶ The entire octree is contained within a cube of scale s_{max}
- ▶ Each children have dimensions that are half the parent's ones and *scale - 1*
- ▶ Each cube is axis-aligned

SVO Traversal

- ▶ We want to know the timestamp at which the ray hits the cube's faces
- ▶ Solving the ray equation $r(t) = o + t * d$ for the x-axis gives

$$t_x(x) = \frac{1}{d_x}x + \frac{-p_x}{d_x}$$


we can do the same for t_y and t_z and precompute the coefficients

ADVANCE

- ▶ Proceed to the next sibling voxel of same dimensions
- ▶ We just need to change idx
- ▶ Each cube is axis-aligned and defined by two opposite vertices (x_0, y_0, z_0) and (x_1, y_1, z_1) such that
 - $t_x(x_0) \leq t_x(x_1)$
 - $t_y(y_0) \leq t_y(y_1)$
 - $t_z(z_0) \leq t_z(z_1)$
- ▶ The t -values span intersected by the cube is given by
 - $tc_{min} = \max(t_x(x_0), t_y(y_0), t_z(z_0))$
 - $tc_{max} = \min(t_x(x_1), t_y(y_1), t_z(z_1))$


ADVANCE

- ▶ In other words there are 6 axis-aligned planes and when the ray enters the first 3 it is inside the cube, when it exits at least one of the last 3 it exits the cube
- ▶ We can determine the next voxel of same scale by comparing $t_x(x_1)$, $t_y(y_1)$, $t_z(z_1)$ with tc_{max} and for each equality we need to flip the corresponding bit in idx



PUSH


- ▶ Proceed to the child voxel that the ray enters first
- ▶ We evaluate t_x , t_y , t_z at the voxel's center (just 3 planes) and compare them against tc_{min}
- ▶ The comparison gives us which bits set in idx




POP

- ▶ Proceed to the next sibling of the highest ancestor that the ray exits
- ▶ We need a POP when an ADVANCE creates an idx that disagree with ray's direction
- ▶ This corresponds with the ray exiting not only from the current node, but also from its parent


POP: Common Ancestor



POP: Common Ancestor



POP: Common Ancestor



Position at 024			
	0	1	2
x	0	0	0
y	0	0	1
z	0	0	0
	0	2	4

Position at 062			
	0	1	2
x	0	0	1
y	0	0	1
z	0	0	0
	0	6	2

POP: Common Ancestor

My implementation differs from the paper's one because I don't need to descend from the highest common ancestor to collect contours' information

Pseudocode

```
INITIALIZE
  1:  $(t_{min}, t_{max}) \leftarrow (0, 1)$ 
  2:  $t' \leftarrow \text{project cube}(root, ray)$ 
  3:  $t \leftarrow \text{intersect}(t, t')$ 
  4:  $h \leftarrow t'_{max}$ 
  5:  $\text{parent} \leftarrow root$ 
  6:  $idx \leftarrow \text{select child}(root, ray, t_{min})$ 
  7:  $(pos, scale) \leftarrow \text{child cube}(root, idx)$ 
  8: while not terminated do
    9:    $tc \leftarrow \text{project cube}(pos, scale, ray)$ 
    10:  if voxel exists and  $t_{min} \leq t_{max}$  then
      11:    if voxel is small enough then return  $t_{min}$ 
      12:     $tv \leftarrow \text{intersect}(tc, t)$ 
      13:    if voxel has a contour then
      14:       $t' \leftarrow \text{project contour}(pos, scale, ray)$ 
      15:       $tv \leftarrow \text{intersect}(tv, t')$ 
      16:    end if
      17:    if  $tv_{min} \leq tv_{max}$  then
      18:      if voxel is a leaf then return  $tv_{min}$ 

INTERSECT
  19: if  $tc_{max} < h$  then  $\text{stack}[scale] \leftarrow (\text{parent}, t_{max})$ 
  20:    $h \leftarrow tc_{max}$ 
  21:    $\text{parent} \leftarrow \text{find child descriptor}(\text{parent}, idx)$ 
  22:    $idx \leftarrow \text{select child}(pos, scale, ray, tv_{min})$ 
  23:    $t \leftarrow tv$ 
  24:    $(pos, scale) \leftarrow \text{child cube}(pos, scale, idx)$ 
  25:   continue
  26: end if
  27: end if
  28:  $oldpos \leftarrow pos$ 
  29:  $(pos, idx) \leftarrow \text{step along ray}(pos, scale, ray)$ 
  30:  $t_{min} \leftarrow tc_{max}$ 
  31: if  $idx$  update disagrees with  $ray$  then
  32:    $scale \leftarrow \text{highest differing bit}(pos, oldpos)$ 
  33:   if  $scale \geq s_{max}$  then return miss
  34:    $(\text{parent}, t_{max}) \leftarrow \text{stack}[scale]$ 
  35:    $pos \leftarrow \text{round position}(pos, scale)$ 
  36:    $idx \leftarrow \text{extract child slot index}(pos, scale)$ 
  37:    $h \leftarrow 0$ 
  38: end if
  39: end while
```

My implementation differs in the INTERSECT and *pos*-related parts

Implementation: Difficulties


- ▶ Difficulties? A lot...
- ▶ Several implementation specific problems that I had to overcome on my own
- ▶ Debugging was really time consuming because there were several factors to take into consideration
- ▶ Even when things (seem to) work, you're afraid that it could be only a lucky case
- ▶ Hofstadter's Law: It always takes longer than you expect, even when you take into account Hofstadter's Law.

But in the end highly rewarding

Repo link:

<https://github.com/Nyriu/RayTracer/tree/octree>

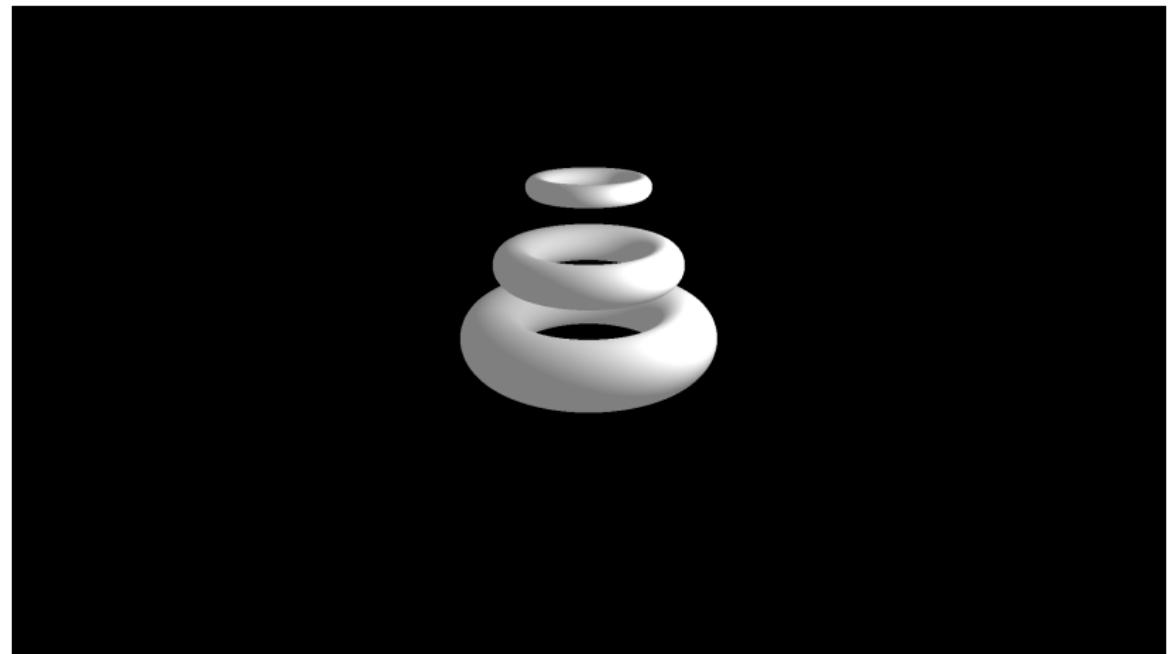
Implementation: Results



1920x1080

SVO	s	μs
No	29	29266361
Yes	31	31826758

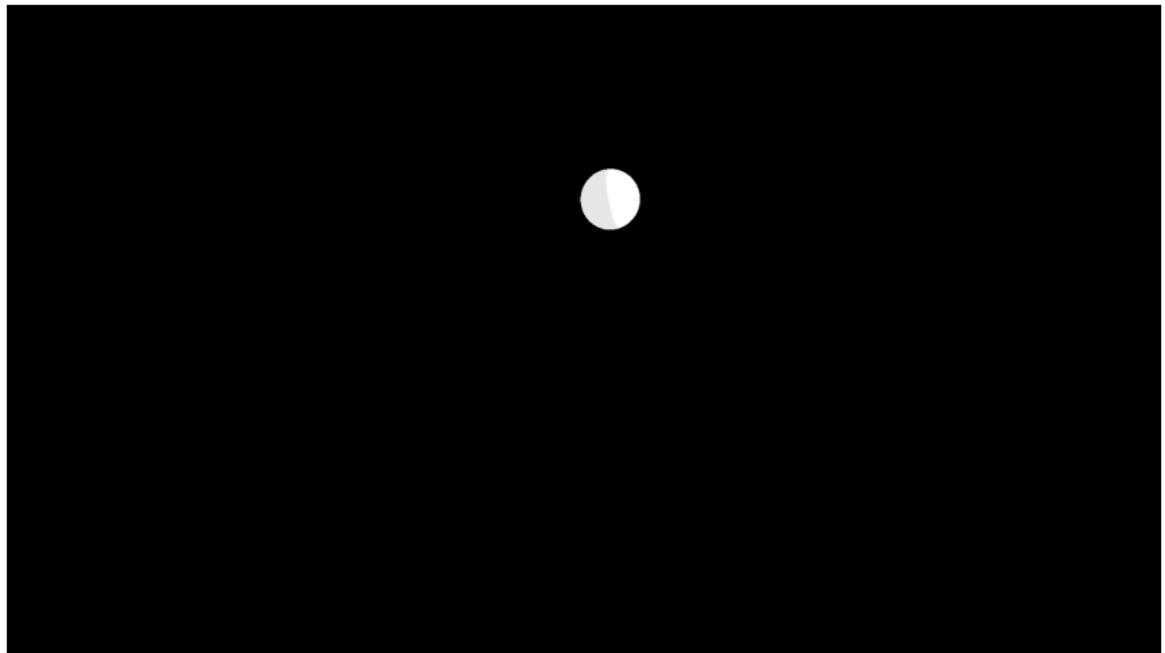
Implementation: Results



1920x1080

SVO	s	μs
No	9	9324068
Yes	9	9558236

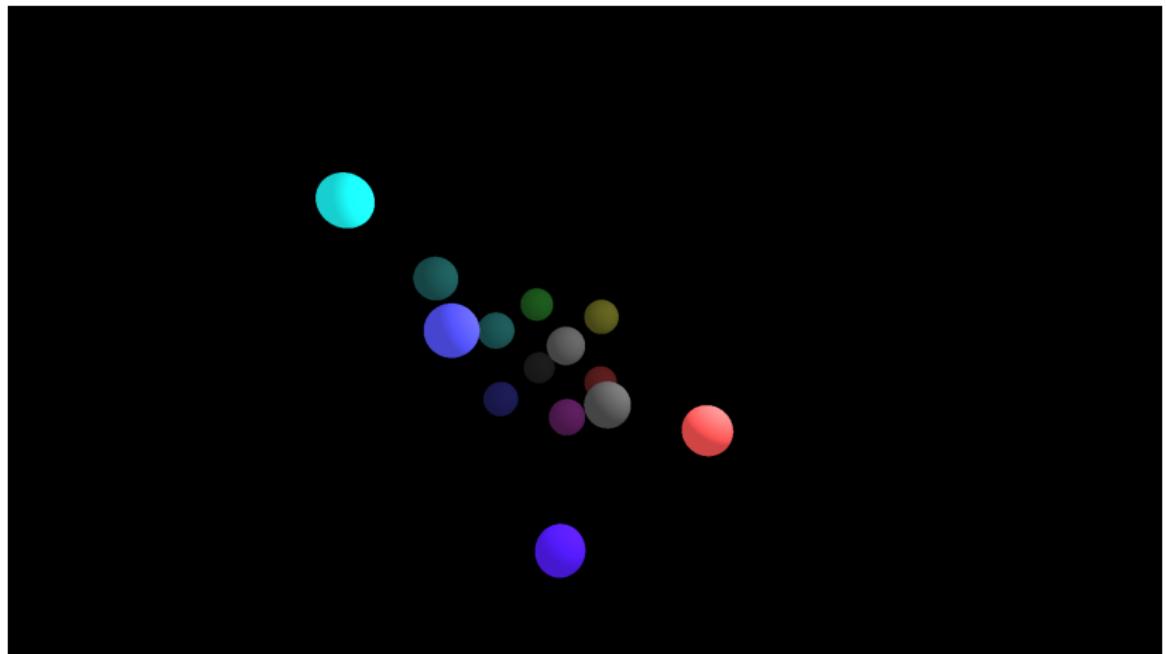
Implementation: Results



1920x1080

SVO	s	μs
No	3	3798347
Yes	2	2344861


Implementation: Results



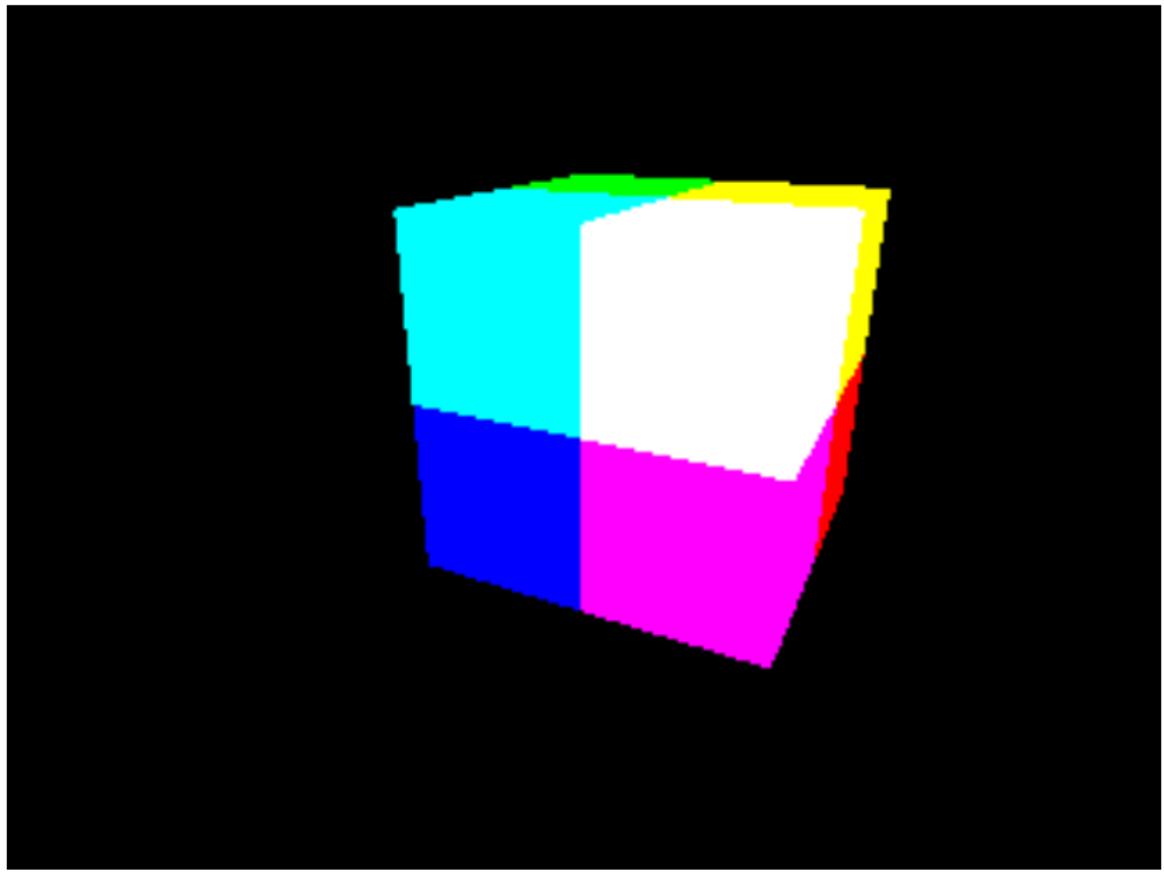
1920x1080

SVO	s	μs
No	31	31014649
Yes	18	18075335


Implementation: Helpers



Implementation: Helpers




Implementation: Initial Bugs



Implementation: Initial Bugs

Implementation: Initial Bugs



Future

- ▶ Cast shadow rays
- ▶ Improve the traversal algorithm and add shortcuts
- ▶ Add materials system
- ▶ Add different sampling patterns for octree generation
- ▶ Parallelism

References I

-  Jhon C. Hart, *Sphere tracing: a geometric method for antialiased ray tracing of implicit surfaces*, The Visual Computer (1996).
-  Inigo Quilez, *Deriving the sdf of a box.*
-  _____, *Distance functions.*
-  Tero Karras Samuli Laine, *Effcient sparse voxel octrees*, NVIDIA Research (2010).

Thank You